



# Manipulating Table Data

CMPE 2400

Databases

# Adding New Records

- ▶ Recall that a record is the same as a row of data
- ▶ Be mindful of primary and foreign key relationships when inserting a record into a table
- ▶ A primary key constraint will reject duplicate column values
- ▶ A foreign key constraint will prohibit entering values into the column that do not already exist in the primary key table

# INSERT INTO x

- ▶ Adding an individual record to a **table** will be accomplished using the following:
  - ▶ **insert into** allows you to choose the **table** that will be affected, and to choose the columns for which you will enter data.
    - ▶ The chosen columns do not need to be listed in the same order as they are listed in the **table**.
    - ▶ If you are going to supply **values** for all columns, you may omit the column list entirely.
    - ▶ If a column has a **default constraint**, then the value for that column can be left out of the column list, and the **default** value will be used in its place.

# Adding New Records

- ▶ Syntax 1- Including Column names

```
insert into Database.Schema.Table  
  (Col1Name, Col2Name, Col3Name, etc)  
values  
  (Value1, Value2, Value3, etc)
```

Example:

```
insert into username_Referees.dbo.Refs  
  (RefID, FirstName, LastName)  
values  
  (2, 'John', 'Mills')
```



# Adding New Records (Without using column names)

- ▶ Adding an individual record to a table will be accomplished using the following:
  - ▶ values allows you to specify the data that you will enter into the chosen columns.
    - ▶ When excluding column names, the data *must* appear in the same order the columns occur in the table.
    - ▶ Columns in a table can be left without a value by making it null, if that is allowed by the constraints
    - ▶ As you might imagine, you *must* supply values for all non-null columns.

# Adding New Records

- ▶ Syntax 2- Excluding Column names

```
insert into Database.Schema.Table  
values (Value1,Value2,Value3, etc)
```

Example:

```
insert into username_Referees.dbo.Refs  
values (1,'Sylvia','Venezia','Master')
```

# Create a sample database

- ▶ Create a database for holding some sample tables
  - ▶ Note: You are not responsible for database / table creation at this time.

```
use master  
go  
  
drop database if exists username_Referees  
go  
  
create database username_Referees  
go  
  
use username_Referees  
go
```

```
drop table if exists Refs
```

```
go
```

```
create table Refs
```

```
(
```

```
    RefID      int  not null
```

```
        constraint PK_Refs primary key clustered,
```

```
    FirstName   varchar(20),
```

```
    LastName    varchar(20),
```

```
    CertLevel   char(6)
```

```
    constraint DF_Certs default 'Master'
```

```
)
```

```
go
```

# Create a sample table

# Add data

- ▶ Now enter the following data
- ▶ Note that one of the first names is currently unknown

First Name	Last Name	Cert Level
Sylvia	Venezia	Master
John	Mills	Master
?	Cooper	Master
Liz	Dun	Senior

# Insert first record

- ▶ Add the first record, providing values for all columns, in the order of their creation.

```
insert into username_Referees.dbo.Refs  
values  
(1, 'Sylvia', 'Venezia', 'Master')
```

## Insert second record

- ▶ Add the second record
- ▶ use the default value for the CertLevel of 'Master'

```
insert into username_Referees.dbo.Refs  
  (RefID, FirstName, LastName)  
values  
  (2, 'John', 'Mills')
```

# Insert third record

- ▶ For the 3rd row, the first name must be null
- ▶ This may also be accomplished by omitting the FirstName column from the column list, and leaving out the associated value in the values list.

```
insert into username_Referees.dbo.Refs  
(RefID, FirstName, LastName)  
values  
(3, null, 'Cooper')
```

# Insert last record

- ▶ For the last record, override the default CertLevel value, and specify the columns out of order.

```
insert into username_Referees.dbo.Refs  
  (FirstName, RefID, CertLevel, LastName)  
values  
  ('Liz', 4, 'Senior', 'Dun')
```



# identity Property

- ▶ The identity property for a primary key column causes the column values to auto-increment whenever a new record is added.
- ▶ An initial value and a step value may be defined when creating the identity column.
  - ▶ Both of these values will default to 1 if not specified.
- ▶ This property may only be applied to an integer type column.
- ▶ The identity property may only be used on a single column in a table.

# identity Property

- ▶ Modify the table creation script for the Refs table so that it appears as below, renaming the table being created as MoreRefs, and changing the names of the constraints.

```
drop table if exists MoreRefs
go

create table MoreRefs
(
    RefID      int identity(10,10) not null
        constraint PK_MoreRefs primary key clustered,
    FirstName   varchar(20),
    LastName    varchar(20),
    CertLevel   char(6)
        constraint DF_MoreCerts default 'Master'
)
go
```

# identity Property Example

- ▶ Add the following record, and then look at the table to see that the **identity** column has been auto-populated.
  - ▶ Try running the query multiple times and see that it doesn't break.


```
insert into username_Referees.dbo.MoreRefs  
    (FirstName, LastName, CertLevel)  
values  
    ('Harry', 'Potter', 'Wizard')
```

# Obtaining the value of an identity property

- ▶ When we insert data into a table containing an identity column, the value generated for the identity column will also be copied into the variable `@@identity`.
- ▶ We can then use this value for insertion into other tables



# Adding Multiple Records

- ▶ Multiple records may be added with a single insert into statement by using a select query **in place of** the values clause of the insert into structure
  - ▶ The number of columns in the insert into statement and the select query must be the same
  - ▶ The data types must match
  - ▶ The columns in the select query may come from multiple tables through the use of a join(s)
- 

# Adding Multiple Records

## ▶ Syntax

```
insert into Database.Schema.NewTable  
  (Col1, Col2, Col3, ... )  
    select Col1, Col2, Col3, ...  
      from OldTable(s)  
    [where FilteringExpression]
```

# Inserting Multiple Records Example

- ▶ Using the same `create table` script as for the `MoreRefs` table, change the `table` name to `EvenMoreRefs`, rename the constraints, and run the script.

```
if exists
(
    select      [name]
    from sysobjects
    where      [name] = 'EvenMoreRefs'
)
drop table EvenMoreRefs
go

create table EvenMoreRefs
(
    RefID          int identity(10,10) not null
        constraint PK_EvenMoreRefs primary key clustered,
    FirstName      varchar(20),
    LastName       varchar(20),
    CertLevel      char(6)
        constraint DF_EvenMoreCerts default 'Master'
)
go
```

# IDENTITY

- ▶ IDENTITY is used to have the DBMS automatically create a unique id for a field
- ▶ For special purposes such as data import, IDENTITY can be temporarily disabled:
- ▶ `set identity_insert username_Referees.dbo.EvenMoreRefs on`
  - ▶ <insert statements with explicit values for the ID column>
- ▶ `set identity_insert username_Referees.dbo.EvenMoreRefs off`
- ▶ Remember: only override the IDENTITY behaviour in appropriate circumstances

# Inserting Multiple Records Example

- ▶ Next, using an `insert into` statement with a `select` query, copy all '**Master**' level referees from the `Refs table` into the `EvenMoreRefs table`.
- ▶ Make the RefIDs match the Refs table

```
set identity_insert username_Referees.dbo.EvenMoreRefs on

insert into EvenMoreRefs (RefID, FirstName, LastName, CertLevel)
    select RefID, FirstName, LastName, CertLevel
    from   Refs
    where  CertLevel = 'Master'

set identity_insert username_Referees.dbo.EvenMoreRefs off
```

# Inserting Multiple Records Example

- ▶ Next, get all of the records from the More Refs table, but this time, let the identity column do its job.

```
insert into username_Referees.dbo.EvenMoreRefs  
        (FirstName, LastName, CertLevel)  
select    FirstName, LastName, CertLevel  
from      username_Referees.dbo.MoreRefs
```

# Creating and Filling Tables

- ▶ The `select ... into` statement may be used to create an entire **table** of data, and populate the **table** using columns from another **table(s)**.
- ▶ Because a new **table** is created as part of the `select ... into` process, check to see whether the **table** exists before executing.
- ▶ You may use a `where` clause to filter rows out of the **table** from which you are **selecting** your column.
  - ▶ If you want only the table structure but no row data, then provide a filter that can never be true. ( ie. `ID = null` )
- ▶ You may also use a **join** to source your data from more than one **table**.

# Creating and Filling Tables

```
drop table if exists TargetTable  
go  
  
select Col1, Col2, Col3, ...  
into      TargetDatabase.dbo.TargetTable  
from      SourceDatabase.dbo.SourceTable  
[where     Filter Expression(s)]
```

# Trying it Out

- ▶ Using `select ... into`, `create` a replica of the `EvenMoreRefs table` called `GettingTiredOfRefs`.
  - ▶ You should see that the `table` columns you included in the `select` portion of the statement are all present.
    - ▶ Notice that the properties of the columns are all consistent. This would include the `identity` property if it had been set.
    - ▶ On the other hand, notice that the `constraints` are all **missing**.
  - ▶ If you were to open the `table`, then you would also see that all of the data is present.

# Creating and Filling Tables Example

- ▶ Using `select ... into`, `create` a replica of the `EvenMoreRefs` table called `GettingTiredOfRefs`.


```
drop table if exists GettingTiredOfRefs  
go  
  
select      RefID, FirstName, LastName, CertLevel  
into        username_Referees.dbo.GettingTiredOfRefs  
from        username_Referees.dbo.EvenMoreRefs  
where       CertLevel <> 'Wizard'
```

## Exercise 2

- ▶ Using select Into, create a table called AdditionalEmployees in the username\_DB, we created previously, and make a copy of the EmployeeId, LastName, FirstName and HireDate of all employees from the NorthwindTraders database into this new table.

# Temporary Tables

- ▶ A temporary table may be created when you do not wish the table to persist after you are finished with your current operation.
  - ▶ The table will exist for the duration of your connection to the database.
    - ▶ A procedure is considered to form its own connection to the database. Thus once the procedure ends, any temporary tables created within are automatically destroyed. (We'll talk about stored procedures later in the course)
    - ▶ In order for a table to be viewed as temporary, just preface its name with a hash ( # ) when it is created.



# Temporary Tables

- ▶ Temporary tables may be created using the standard table creation syntax, but at this time we won't be considering it.
- ▶ For the exercises you shall see in this section, only the select ... into situation will benefit from use of temporary tables

# Temporary Tables

- ▶ Consider the previous example using `select .. into`. In order to make the target table temporary, the query should have looked as follows.

```
select RefID, FirstName, LastName, CertLevel  
into #GettingTiredOfRefs  
from username_Referees.dbo.EvenMoreRefs  
where CertLevel <> 'Wizard'
```

# Deleting and Updating Records

- ▶ The **delete** keyword is used to remove entire records from a table.
- ▶ The **update** keyword is used to modify the contents of a column within an existing record.

# Deleting Table Data

- ▶ The delete statement removes zero or more entire records from a table
  - ▶ You may not use delete to remove data from only one column within a table record.
  - ▶ You may not delete records from multiple tables.
- ▶ The following statement will delete **all records** from the table GettingTiredOfRefs
  - ▶ Note that when all records are deleted, that the structure of the table itself remains unchanged.

```
delete from username_Referees.dbo.GettingTiredOfRefs
```

# Deleting Table Data

- ▶ Usually when the delete statement is used, it is focused on data that fits a certain set of parameters.
- ▶ Using the where clause we may set a filter that only deletes records that meet these parameters.

```
delete  
from username_Referees.dbo.EvenMoreRefs  
where CertLevel = 'Master'
```

# Updating Table Data

- ▶ The update statement allows you to change the contents of a record.
- ▶ Syntax:

```
update Database.Schema.TableName  
set     ColumnName1 = NewValue1,  
        ColumnName2 = NewValue2  
[where   FilterExpression]
```

# Updating Table Data

- ▶ Use the where clause to set a filter that only updates records that meet the specified conditions.

```
update    username_Referees.dbo.Refs  
set        FirstName = 'Old Man',  
          LastName = 'Winter'  
where      CertLevel = 'Senior'
```

# Updating Table Data

- ▶ Note that you can even perform an update on column by using the previous value in the column and performing some modification of it. We'll try it in the next exercise

## Exercise3

There are two parts to this exercise:

Using a **select into** statement, create a new table called **Refs2** in your Referees database, by copying all data from **Refs** into this table

Write an **update** statement that replaces all the cert level values in **Refs2** by their corresponding first 3 characters only

# Referential Integrity

- ▶ As you will discover, if you have not already done so, delete may fail if you attempt to remove a record from a primary key column which is referenced as a foreign key in another table.
- ▶ The reason this will fail is that if the delete is allowed from the primary key field, all values referencing the field will be in violation of the associated foreign key constraint.
- ▶ Careful: if the “cascade delete” option was specified, the child records may be automatically deleted instead of receiving an error message



# Referential Integrity

- ▶ The way to solve this problem is to make sure that you code a deletion of the record in the referencing **table** before you attempt to **delete** the **primary key** record.

# Using Transactions

- ▶ Group a number of executions together into transactions to make a block that either all gets saved, or all gets discarded
- ▶ We indicate the beginning of a transaction by a **begin transaction** statement
- ▶ **commit** is used to permanently store the changes to the database
- ▶ **rollback** is used to undo the changes



## Exercise 4

Use a transaction with 2 separate Insert statements to add 2 records to your Refs table. Commit the transaction and do a select \* query to verify that the records have been added.

Use a new transaction to insert 2 new records in the Refs table. In the same transaction include a select \* to verify that the data has been added to the table. Then use a Rollback statement at the end of the transaction to undo the changes.

# Exercise 5

- ▶ Using **Select Into**, create a table in your username\_DB database containing all the product name, unit price and category name of all products from the NorthwindTraders database.